

SCALER – AI Engineer Screening Assignment

The Ask

Build an AI persona of yourself that we can call, chat with, and use to book an interview – end to end, with no human in the loop. If it works under our probing, we'll know you can do the job.

Part A: Voice Agent (35%)

Provide a phone number we can call. Your AI persona should:

- Introduce itself as your AI representative and set context naturally
- Answer questions about your background, skills, and fit for this role
- Handle interruptions, follow-ups, and off-script conversation, no rigid Q&A trees
- Recover gracefully when it doesn't know something (say so, don't invent)
- Ask for our availability, check your real calendar, propose slots, and book a confirmed meeting without human intervention

Stack hints: Vapi · Retell · ElevenLabs · Deepgram · Twilio · LiveKit

Part B: Chat Interface (35%)

Deploy a public chat URL. We should be able to:

- Ask why you're the right person for this role and get a specific, evidence-backed answer
- Ask about any public GitHub repo you've built, it should know the tech stack, purpose, design tradeoffs, and what you'd do differently
- Ask about your resume: education, experience, projects; accurate and specific
- Check availability and book a call directly from chat
- Probe with edge cases, adversarial questions, and prompt injections; it should stay honest and grounded, not hallucinate or break character

The persona must be RAG-grounded over your actual resume and GitHub repos. No hardcoded answers. We may ask questions whose answers exist only in your repo READMEs or commit history – we'll check.

Part C: Evals Report (30%)

Submit a 1–page PDF covering:

- Voice quality: how you measured first-response latency, transcription accuracy, and task completion rate (booking success across N test calls)
- Chat groundedness: your hallucination rate, how you measured it (e.g., judge model, manual labelling, golden Q&A set), and retrieval quality (precision/recall on your own corpus)



- 3 failure modes you discovered, root cause for each, and the fix
- One tradeoff you consciously made (cost vs latency, accuracy vs coverage, etc.) and why
- What you'd build with 2 more weeks

Hard Requirements

Requirement	Detail
Live at submission	Voice and chat must be live when you submit. No screenshots, no videos of past runs.
Real calendar booking	Actual availability and confirmation via Calendly / Cal.com / Google Calendar
RAG grounded	Persona reads your real resume and GitHub : no hardcoded strings
Voice latency	< 2s first response; handles barge-in/interruptions without crashing
Public GitHub repo	Clean README, architecture diagram, setup instructions, and a brief cost breakdown (per call / per chat session)
Eval report	1-page PDF with the measurements above
Loom walkthrough	≤ 4 minutes, covering architecture and one hard problem you solved

What We're Evaluating

System design under real-world constraints, latency engineering, RAG quality, eval rigour, and honesty under pressure. We are not evaluating UI polish.

How to Submit

Fill out the submission form: <https://forms.gle/MrZMGCKikHaFkA3J9>

You'll be asked for:

1. Full name and email
2. Voice agent phone number
3. Public chat URL
4. GitHub repo link
5. Eval report (PDF upload)
6. Loom walkthrough link (≤ 3 min)
7. Total build time (honest estimate - this doesn't affect scoring)

Keep everything live for at least 7 days after submission. We will call and chat unannounced.

